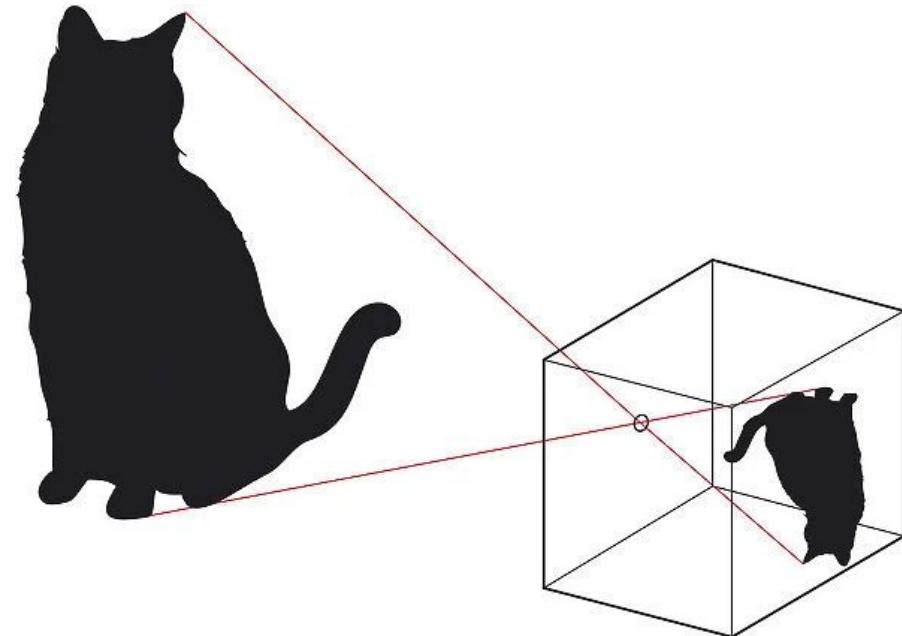


CS-GY 6643

Computer Vision

Lecture 2: Image formation and filtering

Prof. Erdem Varol



Today's menu

Announcements

Image representations (Szeliski 2.2,2.3)

Point operations (Szeliski 3.1)

Intro to image filtering (Szeliski 3.2)

Coding lab



Announcements



Homework 1 out - Due Sep 19 10:59am (1 week)

4 questions ($25 + 25 + 25 + 25 = 100$ points)

6% of total course grade

Please submit a write up in pdf format on brightspace.

No paper submissions.

You can either write on the assignment sheet and scan

...or provide a separate write up (Latex is preferred).

Must submit individually.

Honor code - OK to work with other students, as long as it's stated.

The whole assignment should take <1.5 hrs.

1 point will be deducted for every hour the assignment is late (rounded down).

CS GY 6643 - Computer Vision, Fall 2024

Erdem Varol

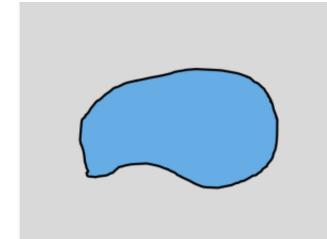
Homework 1

Out: 09/12/2024

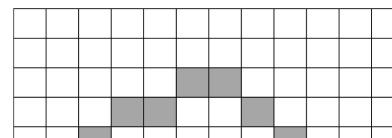
Due: 09/18/2024 11:59 PM

1. Connectivity: Connected component labeling

The famous theorem of Jordan states that every simple closed curve in the plane separates the complement into two connected nonempty sets: an interior region and an exterior. Although this looks simple and intuitive, the proof can get very complex (for those who want to look it up: [Link](#)). Elements are a closed curve and the complement, which is the plane, which is divided into 2 regions. Now let us explore how the continuous case transfers to our discrete domain.



Given the following pixel image with white areas and a gray closed pixel curve. How many connected components can you detect?



Project 1 will be out tomorrow

Due October 4, 10:59am (3 weeks)

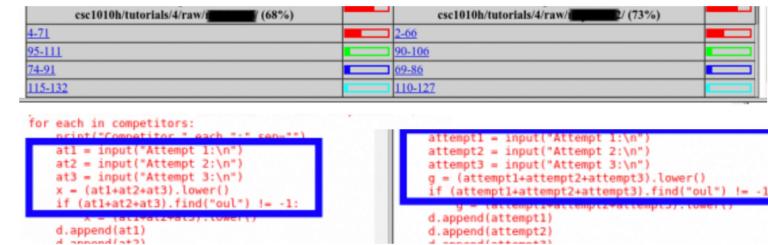
15 % of total course grade.

A python notebook with executed outputs + a pdf write up of each step explained to be submitted on brightspace.

You are encouraged to discuss with other students and look online for examples.

Project must be submitted individually (**We will check for blatant code copying**).

1 point will be deducted for every hour the project is late (rounded down).



The image shows two side-by-side screenshots of Python notebooks. Each notebook has a header bar with the path 'csc1010h/tutorials/4/raw/' and a progress percentage (68% and 73% respectively). Below the header are two tables of student scores:

Score Range	Count
4-71	1
75-111	1
74-91	1
115-132	1

Score Range	Count
2-60	1
90-106	1
69-86	1
110-127	1

Below the score tables are two code snippets highlighted with blue boxes:

```
for each in competitors:  
    print("Competitor", each, ":", sep="")  
    at1 = input("Attempt 1:\n")  
    at2 = input("Attempt 2:\n")  
    at3 = input("Attempt 3:\n")  
    x = (at1+at2+at3).lower()  
    if (at1+at2+at3).find("out") != -1:  
        x = (at1+at2+at3).lower()  
    d.append(at1)  
    d.append(at2)  
    d.append(at3)
```

```
attempt1 = input("Attempt 1:\n")  
attempt2 = input("Attempt 2:\n")  
attempt3 = input("Attempt 3:\n")  
y = (attempt1+attempt2+attempt3).lower()  
if (attempt1+attempt2+attempt3).find("out") != -1:  
    y = (attempt1+attempt2+attempt3).lower()  
d.append(attempt1)  
d.append(attempt2)  
d.append(attempt3)
```



Final project team formation as a homework

Due September 22, 10:59am (10 days)

Worth 2% of total course grade

Each student must submit the names and NYU ID's of teammates

Also, each student must submit the tentative final project area.

To be done brightspace.

Use the match-making spreadsheet to declare project teams + ask to join other teams.

<https://docs.google.com/spreadsheets/d/1RkjD30T0BsH8VO79zaxYJUE6GJFWTgVNu7nE4SJ1UYw/edit?usp=sharing>

	A	B	C	D	E	F
1	Project area	Project tentative title	Team member 1 name	Team member 1 email	Team member 2 name	Team member 2 email
2	Image segmentation	TBA	Firstname Lastname 1	f1234@nyu.edu	Firstname Lastname 2	f1235@nyu.edu
3	Object tracking	Cell tracking in C. elegans microscopy videos	Firstname Lastname 5	f1238@nyu.edu	ASSIGN ME TEAMMATES	
4						
5						



Project ideas

Check slack channel #project-ideas for suggestions.

Kaggle (and other CV competitions) is a great resource for well benchmarked problems.

The screenshot shows the Kaggle interface. On the left is a sidebar with navigation links: Create, Home, Competitions (which is selected), Datasets, Models, Code, Discussions, Learn, and More. The main area is titled "Computer Vision X". It displays a list of competitions:

- RSNA 2024 Lumbar Spine Degenerative Classification**
Classify lumbar spine degenerative conditions
Featured · Code Competition · 1451 Teams · A month to go
- Kannada MNIST**
MNIST like dataset for Kannada handwritten digits
Playground · Code Competition · 1212 Teams · 5 years ago
- HubMAP + HPA - Hacking the Human Body**
Segment multi-organ functional tissue units
Research · Code Competition · 1174 Teams · 2 years ago
- SenNet + HOA - Hacking the Human Vasculature in 3D**
Segment vasculature in 3D scans of human kidney
Research · Code Competition · 1149 Teams · 7 months ago
- RSNA 2023 Abdominal Trauma Detection**
Detect and classify traumatic abdominal injuries
Featured · Code Competition · 1125 Teams · A year ago
- TReNDS Neuroimaging**
Multiscanner normative age and assessments prediction with brain function, structure, and connectivity
Research · 1047 Teams · 4 years ago
- Google Research - Identify Contrails to Reduce Global Warming**
Train ML models to identify contrails in satellite images and help prevent their formation
Research · Code Competition · 954 Teams · A year ago



Take advantage of office hours to discuss assignments + final project ideas



Instructor: Prof. Erdem Varol

Email: ev2240@nyu.edu

Office Hours: On Zoom, 15 min appointments.



TA: Rishabh Raj

Email: rr4574@nyu.edu

Office Hours: Tuesdays 1-2pm on Zoom.



TA: Malhar Patel

Email: mkp6112@nyu.edu

Office Hours: Mondays 1-2pm on Zoom.

Also in person in 370 Jay Street, room 1166

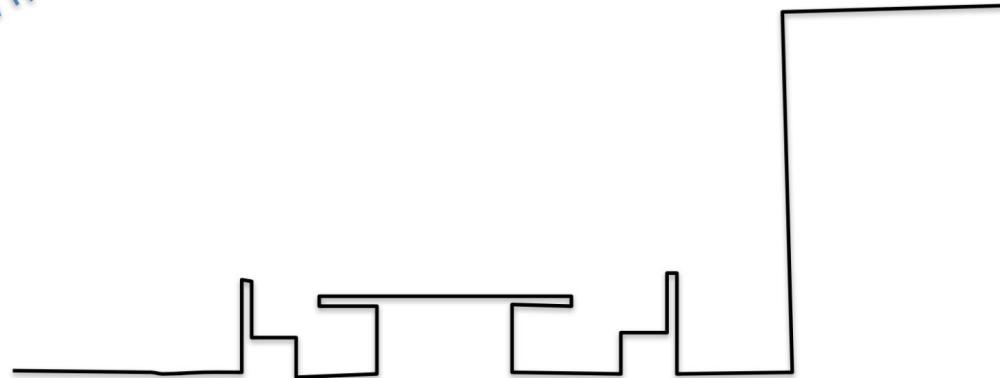
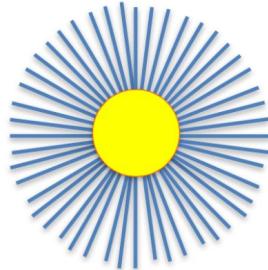


Image formation

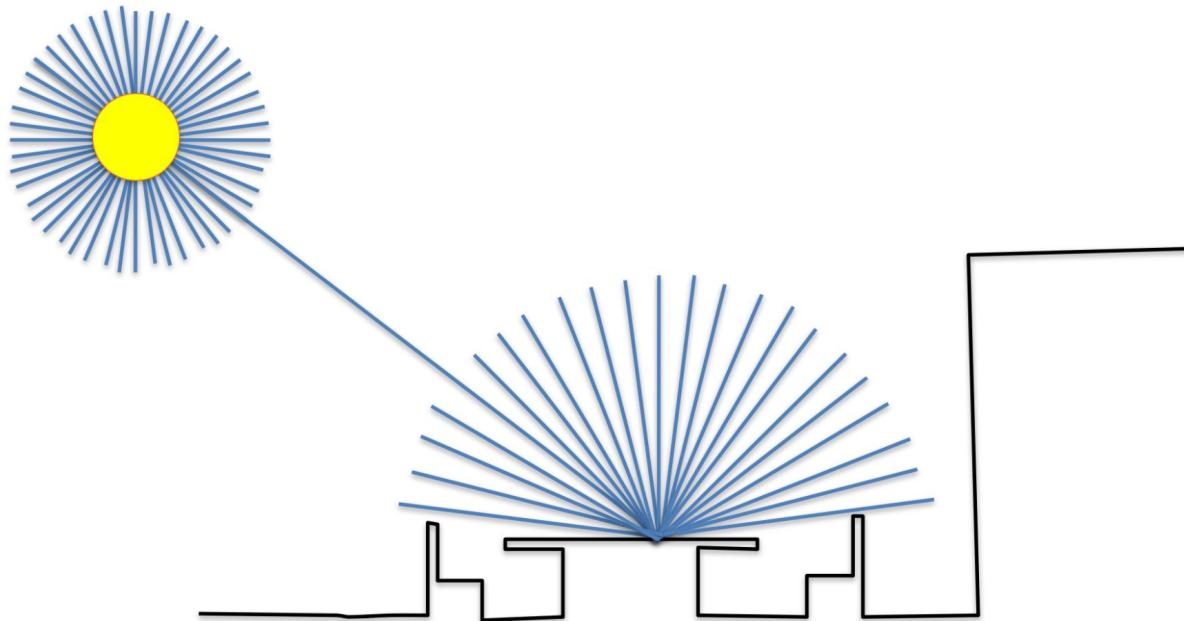
Let there be light!



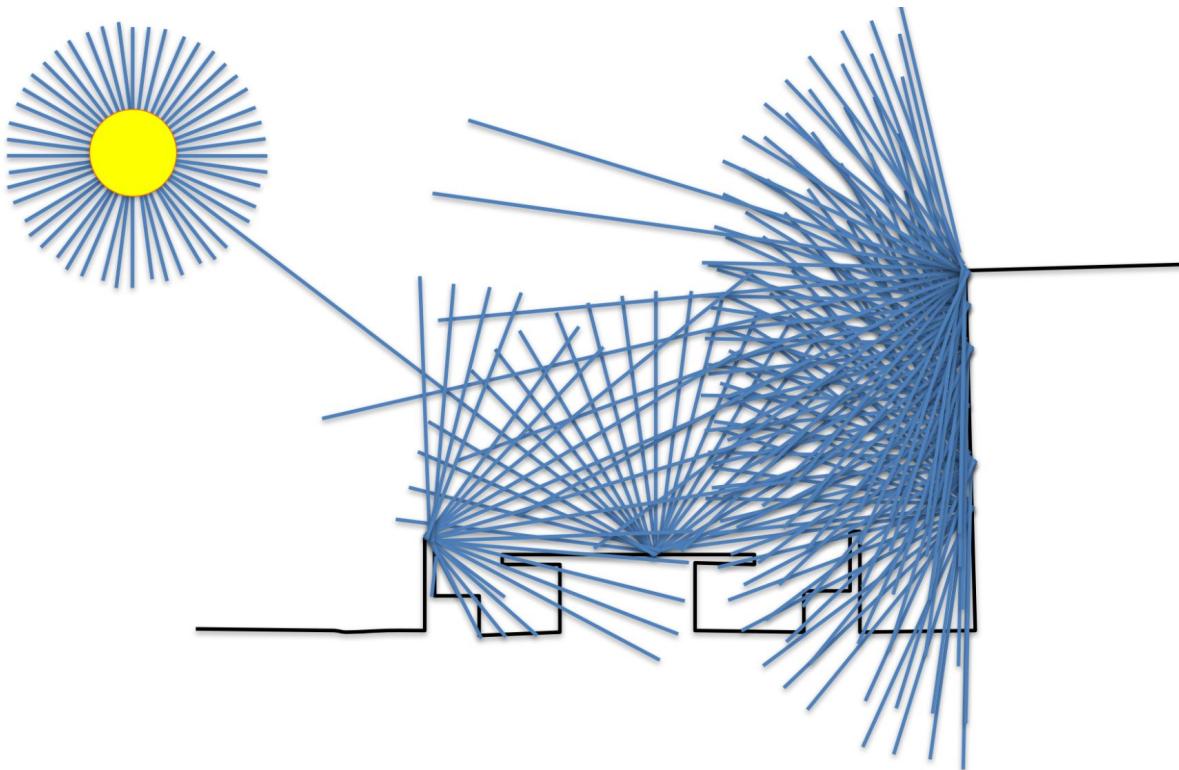
Let there be objects!



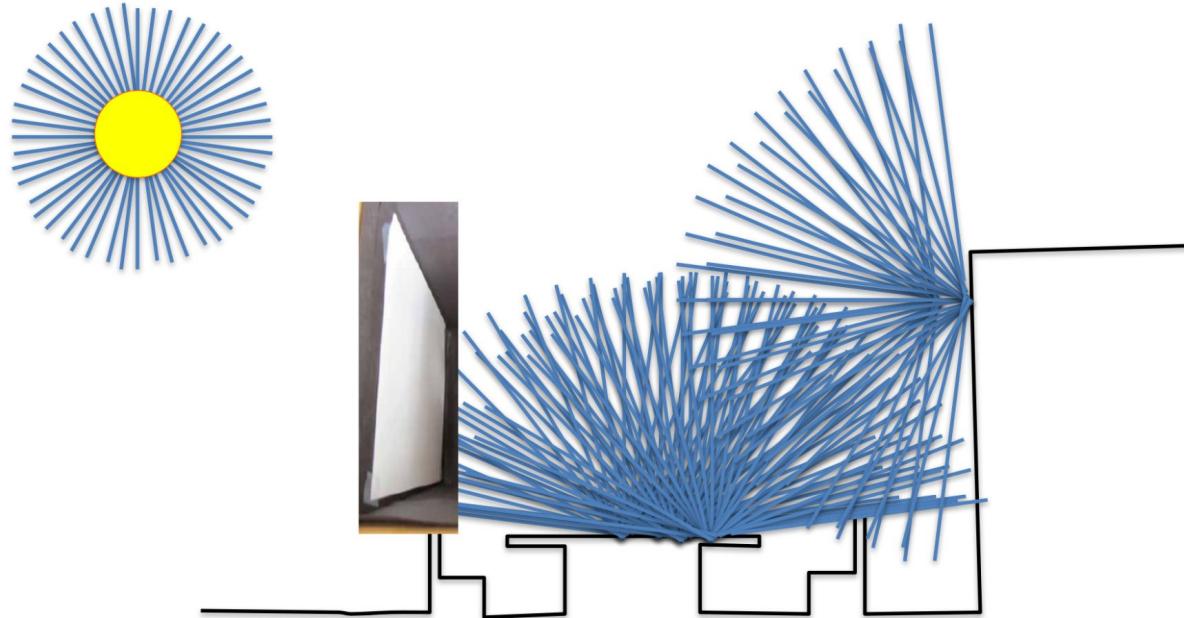
Let there be physics!



Let there be more physics...

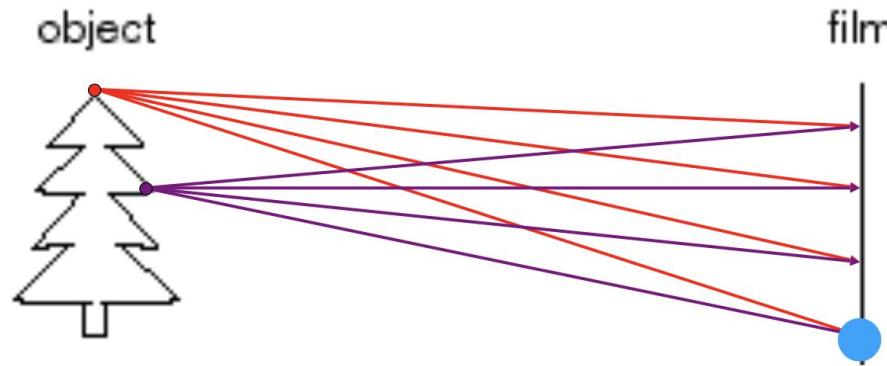


If we put a sensor, what happens?



There will be no image. Why?

Why doesn't this generate an image?



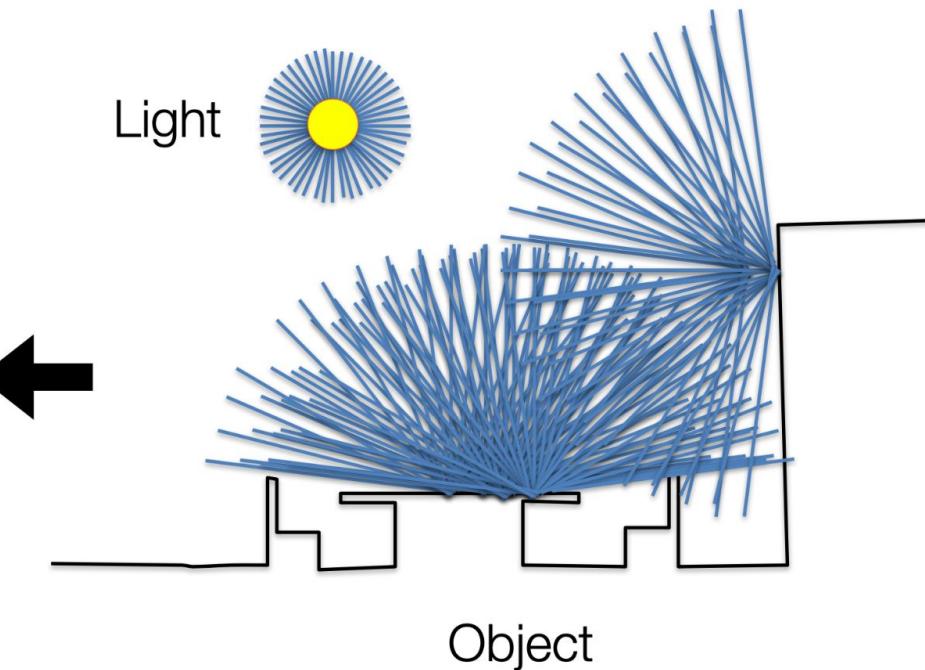
For each point on the film

- it receives light from different directions
- in the end, the color is averaged out

What is needed? A tool to restrict light that comes only from where we want to see...



Camera



Object

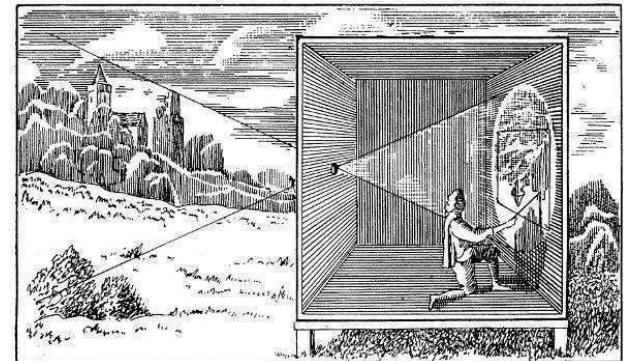
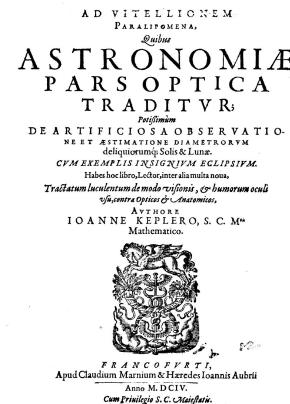
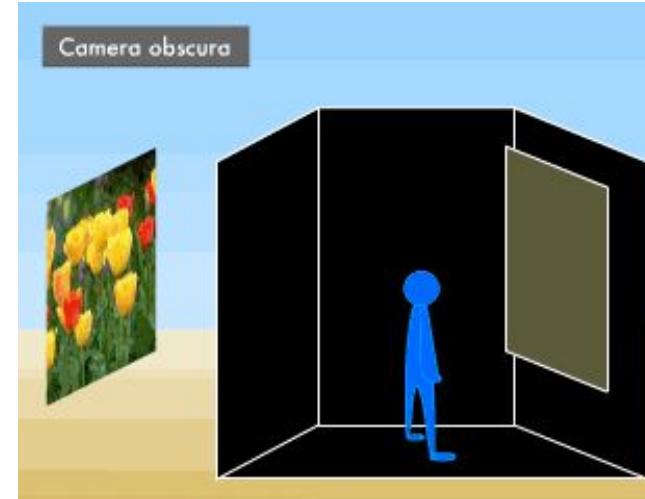
Camera obscura

First camera in history

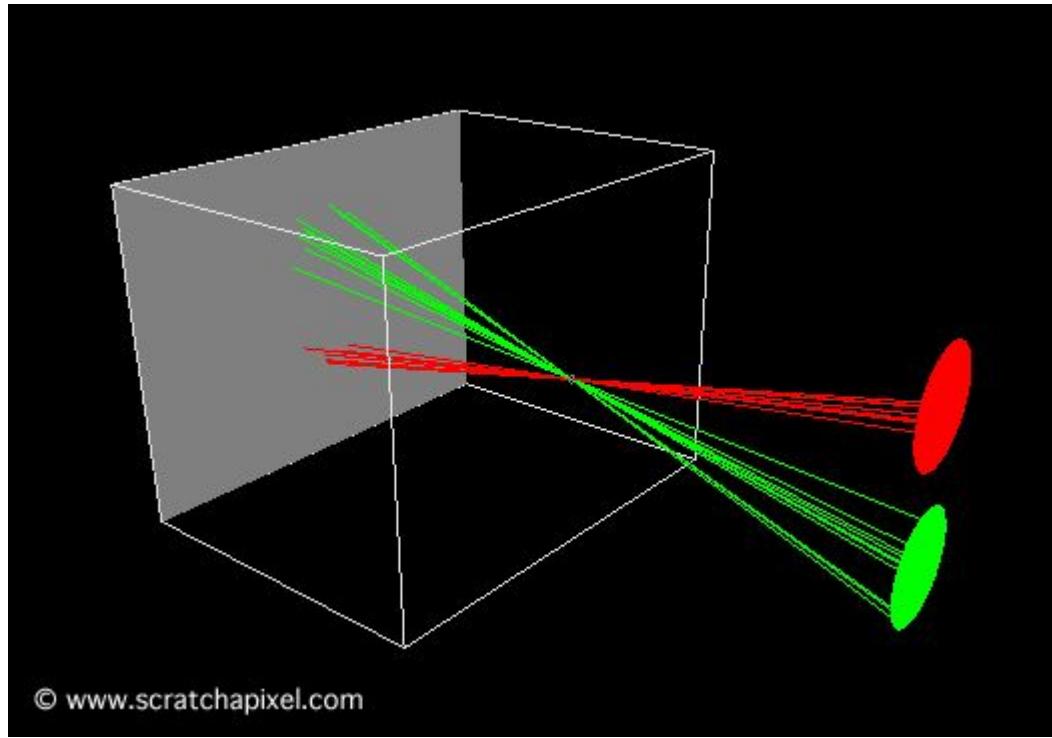
Mentioned in ~500 BC in Chinese writings

Term coined and formalized in 1604 by Kepler

17th century Dutch painters used it for perspective studies.



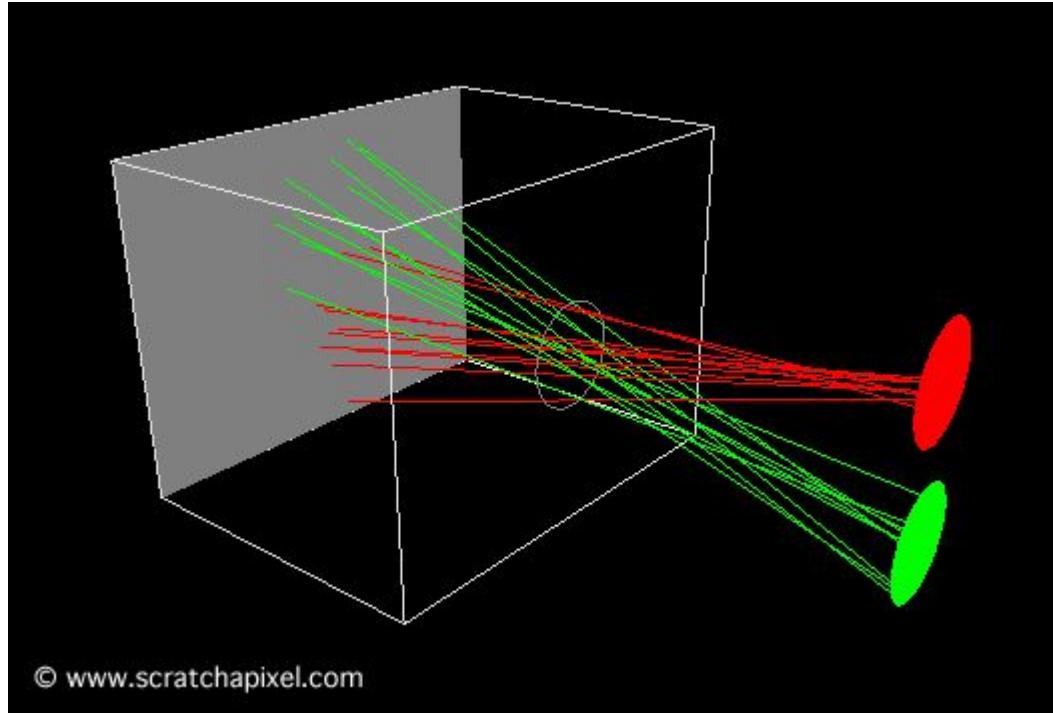
How does it work?



© www.scratchapixel.com



What if the pinhole was bigger?



Same issue as having no camera arises - lots of uncontrolled photons hitting the sensor.

21st century DIY pinhole camera



View outside the class

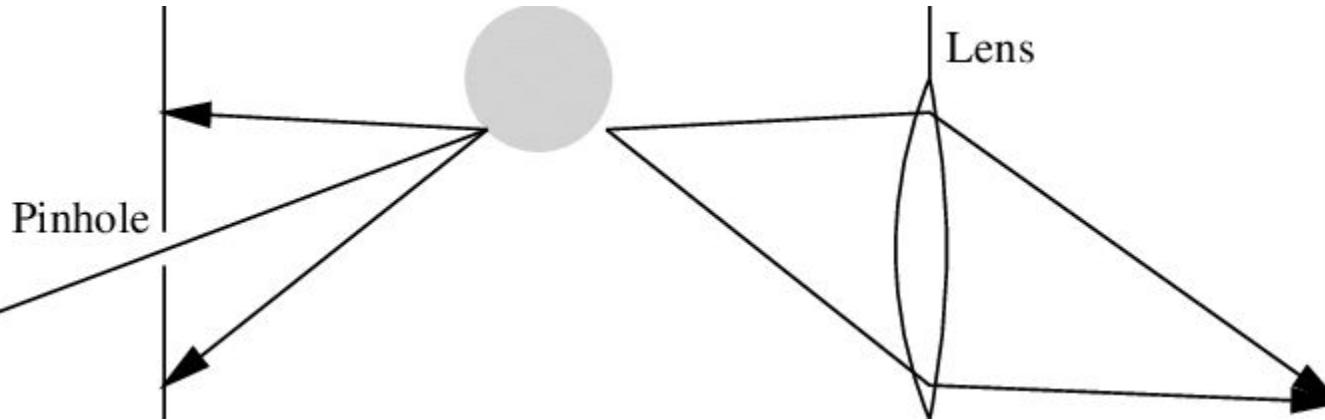


Camera lens



Camera lens replaced by a very small hole

Pinhole vs. an actual lensed camera

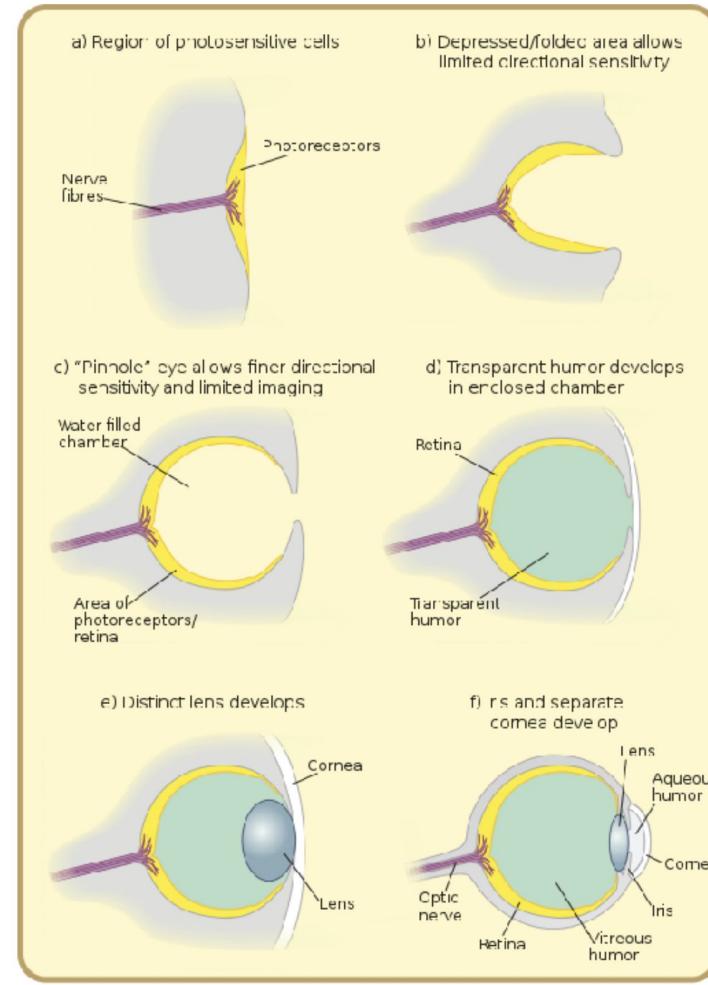


Lens allows for a larger set of photons to be focused on a single point.
Pinhole merely restricts the photons that are scattered in all directions.

Recall first lecture

Primitive “eyes” were pinhole cameras.

Advanced organisms evolved to have lenses in their “cameras”.



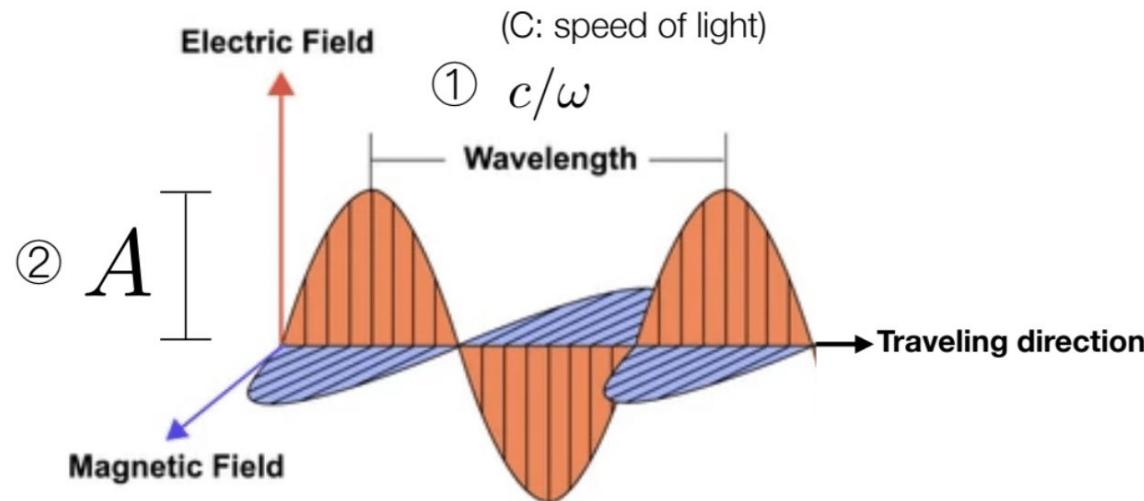
Takeaway:

An image is a focused beam of photons that originate from a specific location in space.

Blur is a collection of photons that come from many random places in space.



What are photons? A type of electromagnetic radiation



① ω : Frequency (how fast it vibrates)

② A : Amplitude (how big it vibrates)

Photons have two qualities



Input image

Brightness



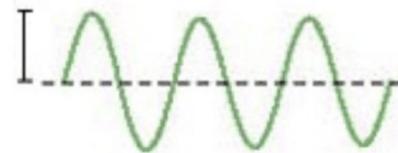
Color



Photon amplitude = brightness

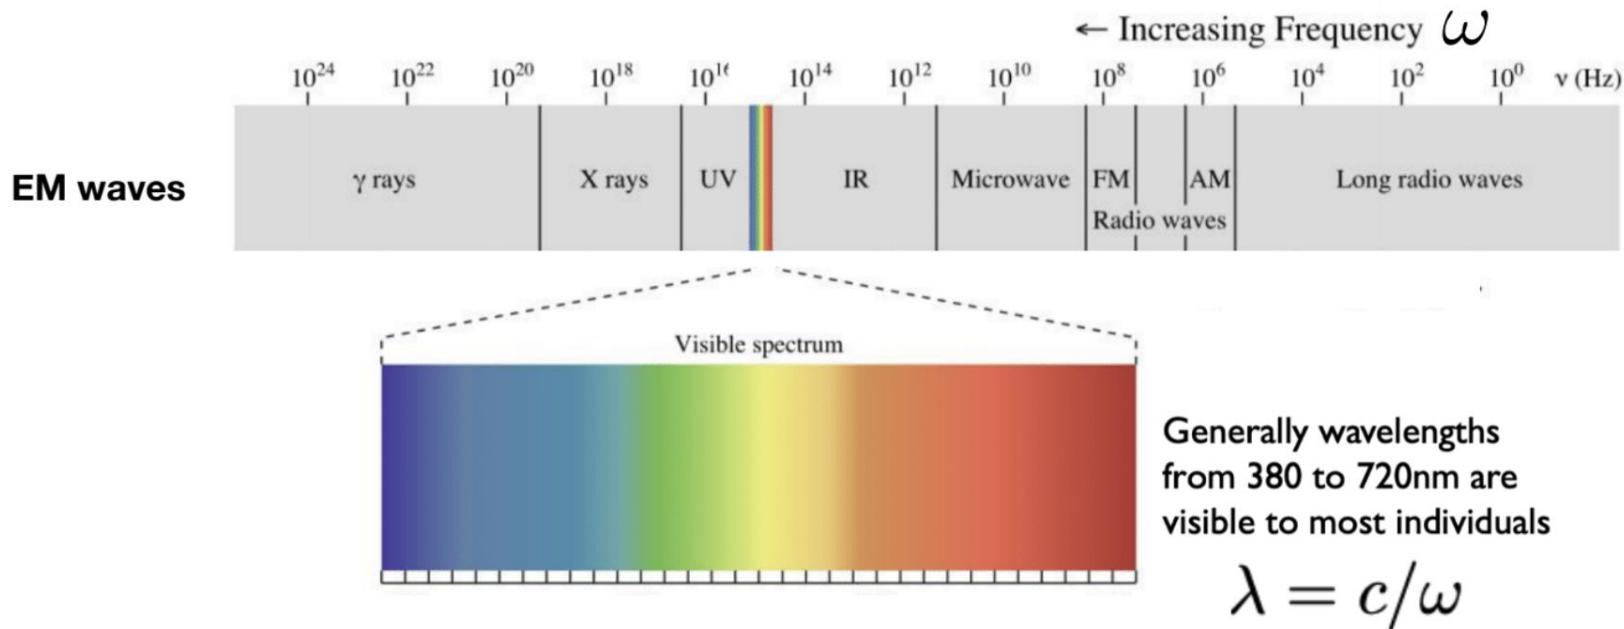


Low brightness



High brightness

Photon frequency = color

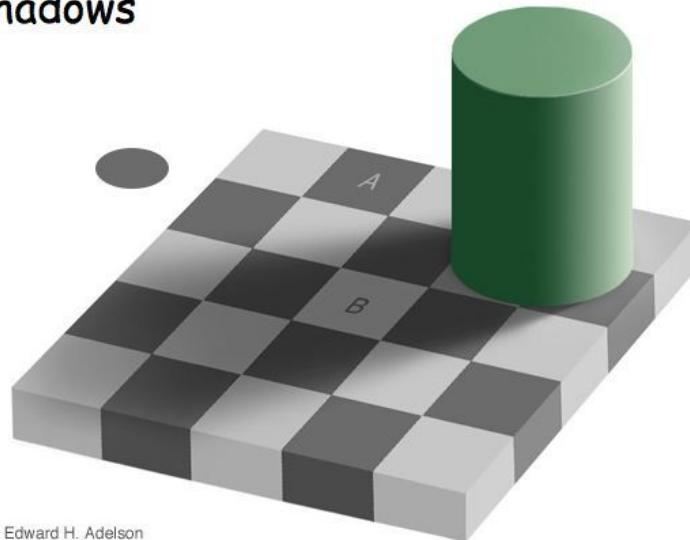


**So far, we've talked about vision from a
purely objective perspective: photons,
brightness, color...**

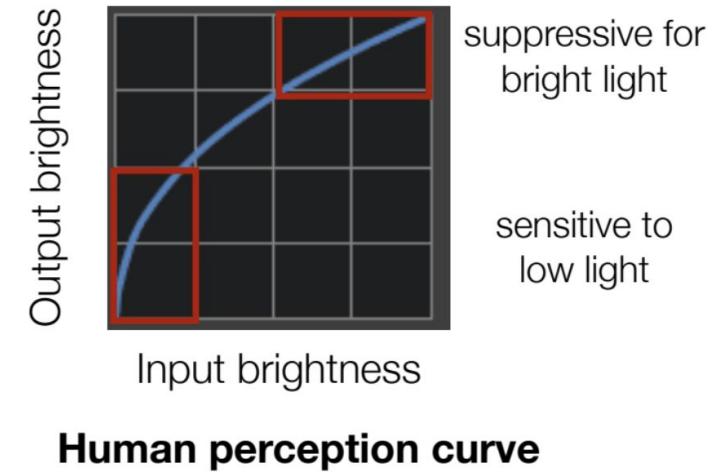
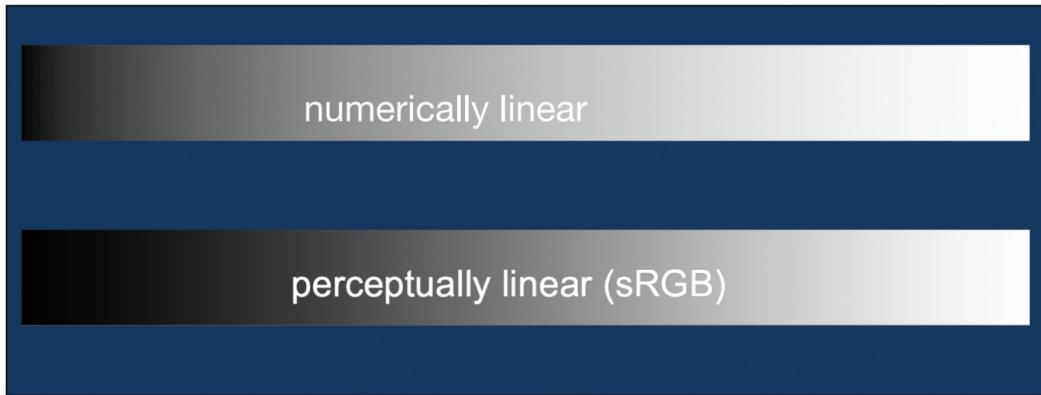


Vision is also subjective = perception

Shadows

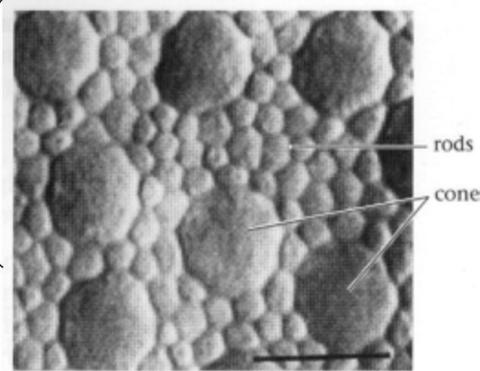
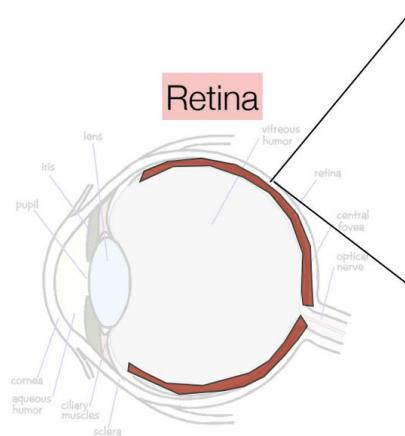


Perceptual brightness is non-linear

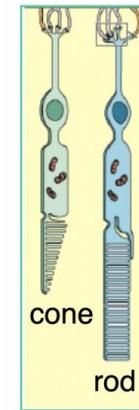


Why? Answer = evolution

Retina: cones and rods (light sensors)



2D retina surface



Depth view
(Into the 2D surface)

Two types of sensors: one for color, one for brightness

Cone

- Specialized in color



Normal cones



Impaired cones
(Color blind)

Rod

- Sensitive to brightness
- gray-scale

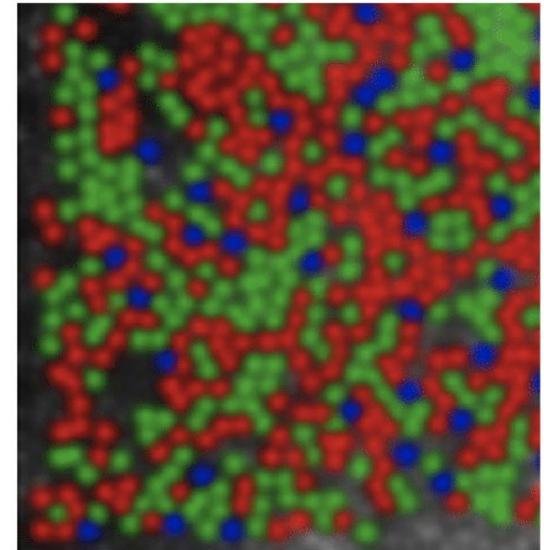
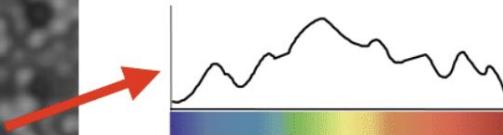
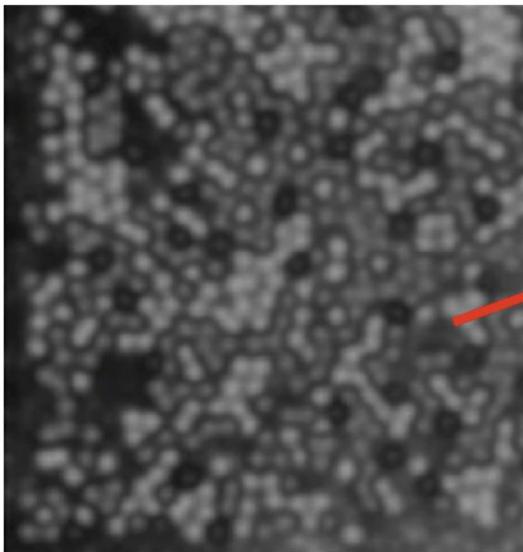


Normal rods



Impaired rods
(Night blindness)

3 types of cones = R, G, B



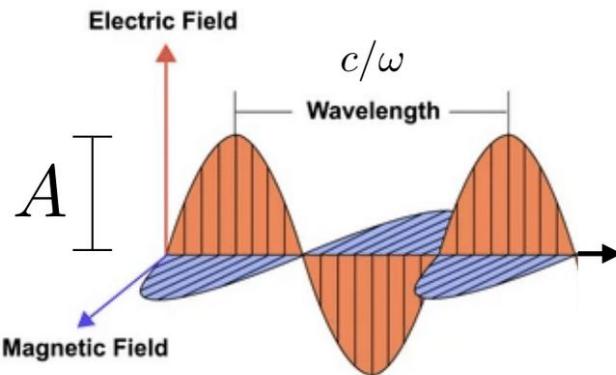
Note: Each cone doesn't pick a single "color" but a range of colors.

Objective vs. subjective vision

Physics (Optics)

- Wave frequency → Color
- Wave amplitude → Brightness

We can measure these objectively with sensors and cameras.



Objective vs. subjective vision

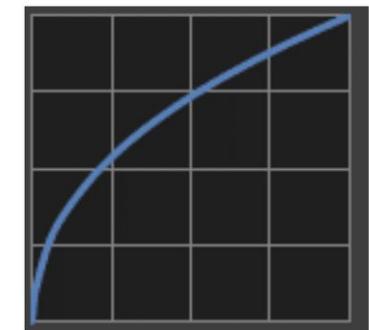
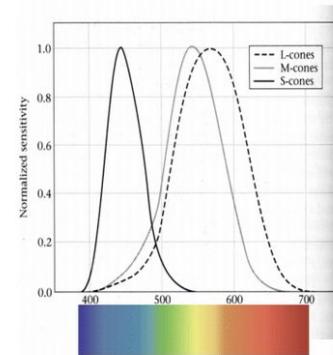
Perception

Color
Brightness

- 3 types of cones
- Non-linear response (**Gamma curve**)

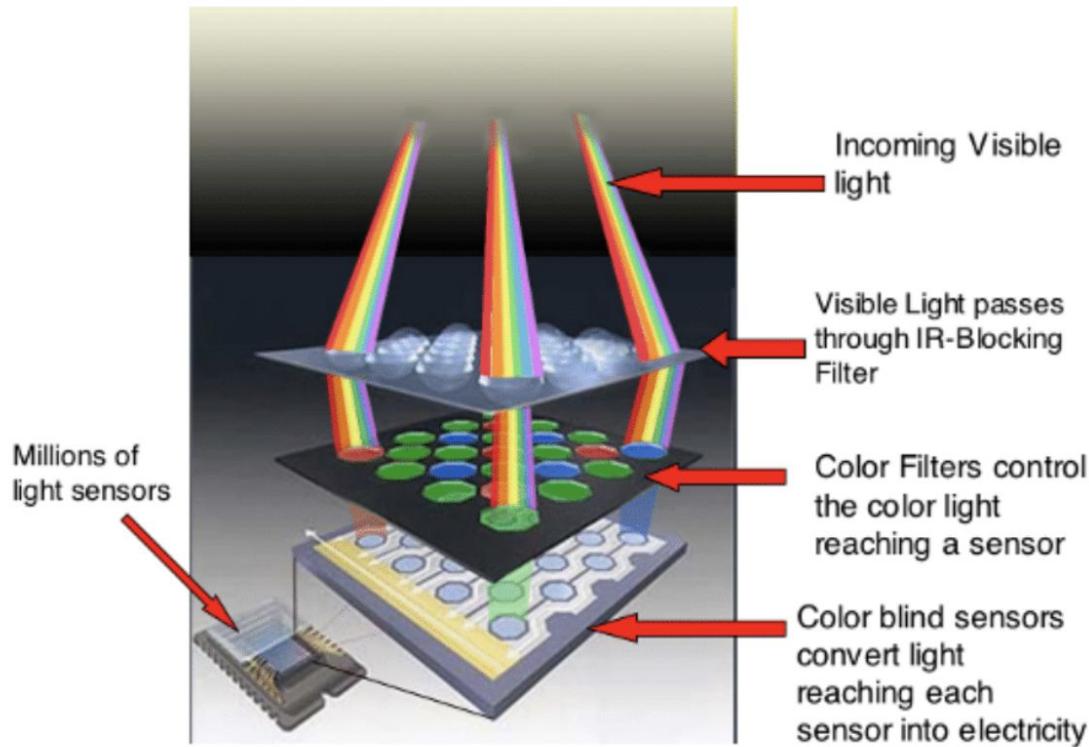
We can't measure these objectively
- instead, we build surrogate
systems to mimic it.

This is computer vision.



How do we build artificial perception systems?

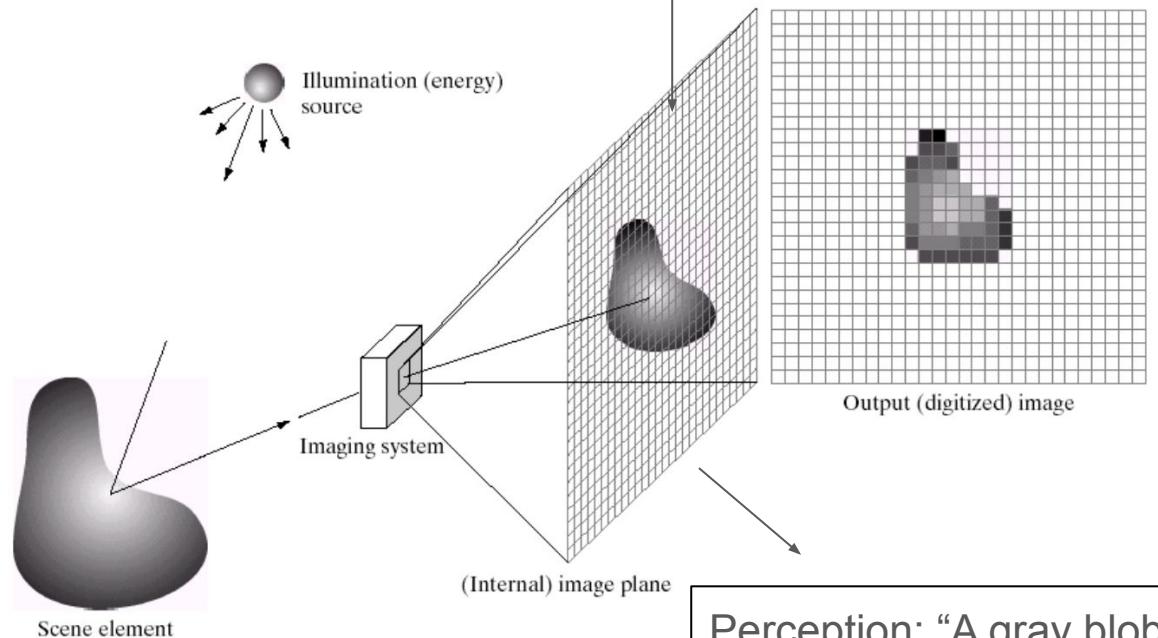
Camera sensors



Digitized signal

Digital images

Each pixel is a sampling of the number of photons in particular spatial bin.



What computers see

243	239	240	225	206	185	188	218	211	206	216	225
242	239	218	110	67	31	34	152	213	206	208	221
243	242	123	58	94	82	132	77	108	208	208	215
235	217	115	212	243	236	247	139	91	209	208	211
233	208	131	222	219	226	196	114	74	208	213	214
232	217	131	116	77	150	69	56	52	201	228	223
232	232	182	186	184	179	159	123	93	232	235	235
232	236	201	154	216	133	129	81	175	252	241	240
235	238	230	128	172	138	65	63	234	249	241	245
237	236	247	143	59	78	10	94	255	248	247	251
234	237	245	193	55	33	115	144	213	255	253	251
248	245	161	128	149	109	138	65	47	156	239	255
190	107	39	102	94	73	114	58	17	7	51	137
23	32	33	148	168	203	179	43	27	17	12	8
17	26	12	160	255	255	109	22	26	19	35	24

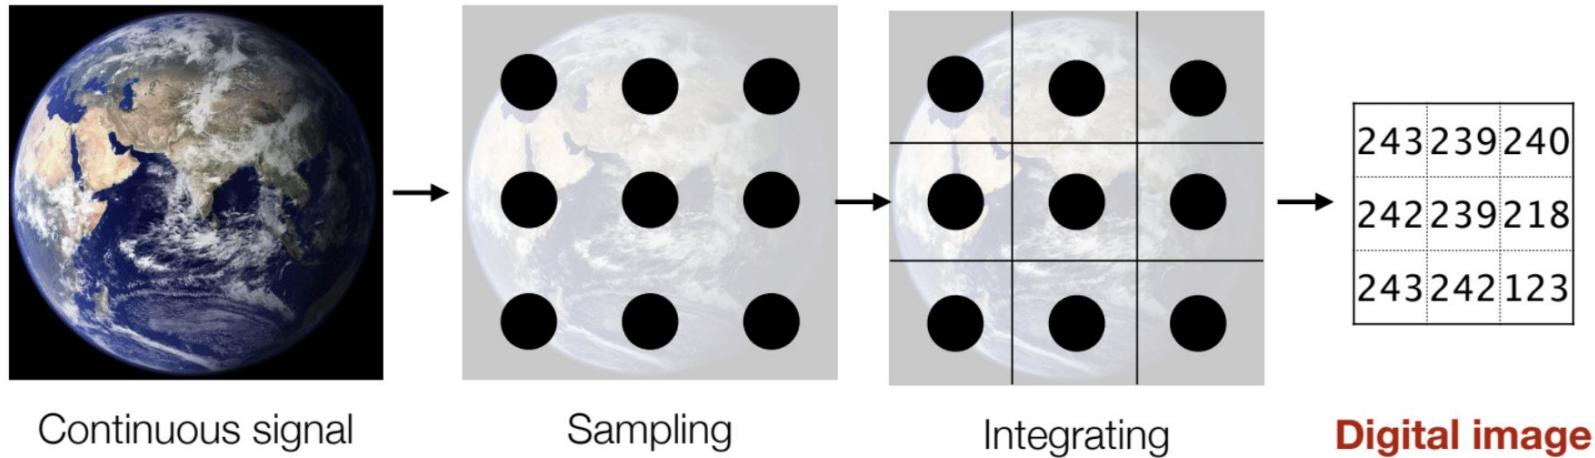


What humans see

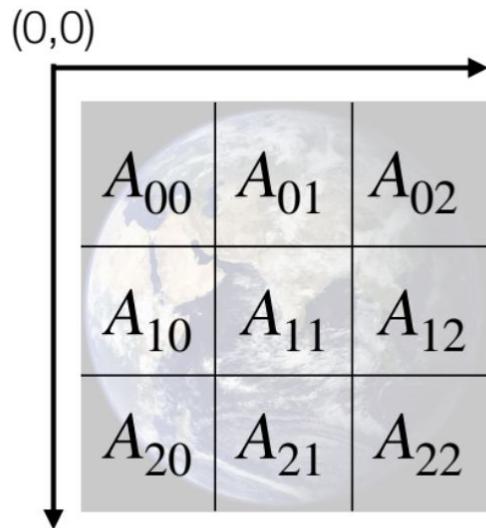
243	239	240	225	206	185	188	218	211	206	216	225
242	239	218	110	67	31	34	152	213	206	208	221
243	242	123	58	94	82	132	77	108	208	208	215
235	217	115	212	243	236	247	139	91	209	208	211
233	208	131	222	219	226	196	114	74	208	213	214
232	217	131	116	77	150	69	56	52	201	228	223
232	232	182	186	184	179	159	123	93	232	235	235
232	236	201	154	216	133	129	81	175	252	241	240
235	238	230	128	172	138	65	63	234	249	241	245
237	236	247	143	59	78		94	255	248	247	251
234	237	245	193	55	33	115	144	213	255	253	251
248	245	161	128	149	109	138	65	47	156	239	255
190	107	39	102	94	73	114	58	17		51	137
23	32	33	148	168	203	179	43	27	17		
17	26		160	255	255	109	22	26	19	35	24



An digital image is a matrix of pixel values



Each pixel has a coordinate



Put matrix on top image

```
A = load_image(..)  
print(A[2,1])
```

243	239	240
242	239	218
243	211	123

In Python



Color images are 3D matrices = tensors



Grayscale image: 2D matrix ($M \times N$)

243	239	240
242	239	218
243	242	123

$A[\text{row}, \text{col}]$



Color image: 3D tensor ($M \times N \times 3$)

243	239	240	
243	239	240	18
243	239	240	18
242	239	218	23
243	242	123	23

$A[\text{row}, \text{col}, \text{chan}]$

52

What makes an image informative

Example matrix: **(uint8)^{3×3}**

243	239	240
242	239	218
243	242	123

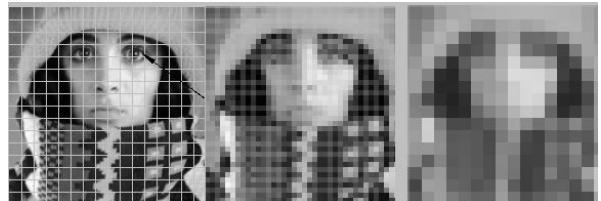
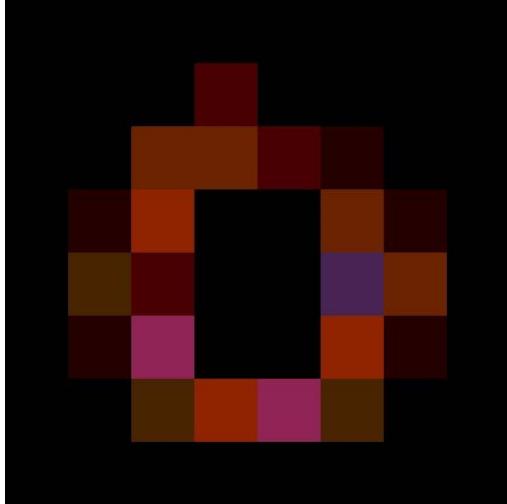
- Bit depth: uint8 → uint16
- Resolution: $3 \times 3 \rightarrow 30 \times 30$
- Color: $30 \times 30 \rightarrow 30 \times 30 \times 3$

Data type uint8: 0-255



Quantization

How pixels are quantized affects image resolution.



Summary I: Digital image = (data type) $N_1 \times N_2 \times \dots \times c$



World



243	239	240
242	239	218
243	242	123

Digital image



data type : uint8, uint16 (**bit depth**)

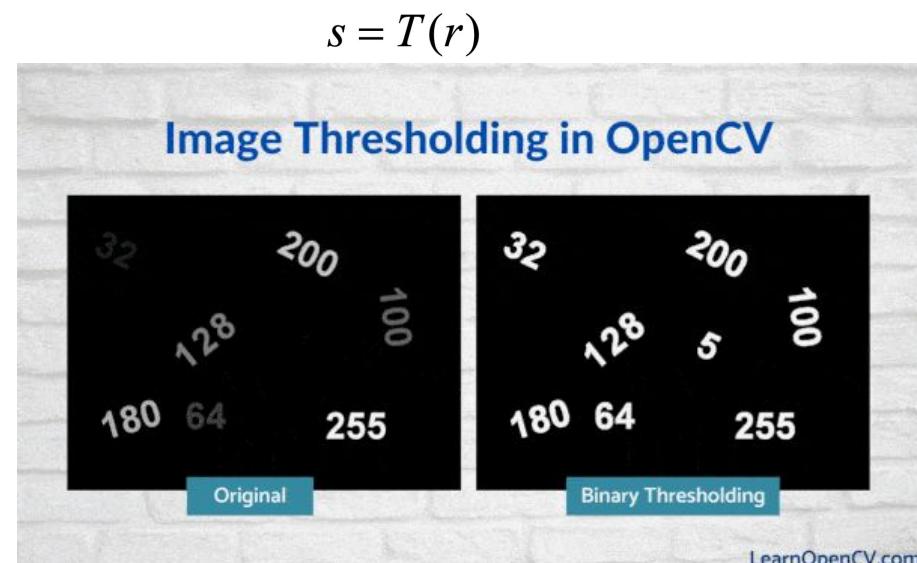
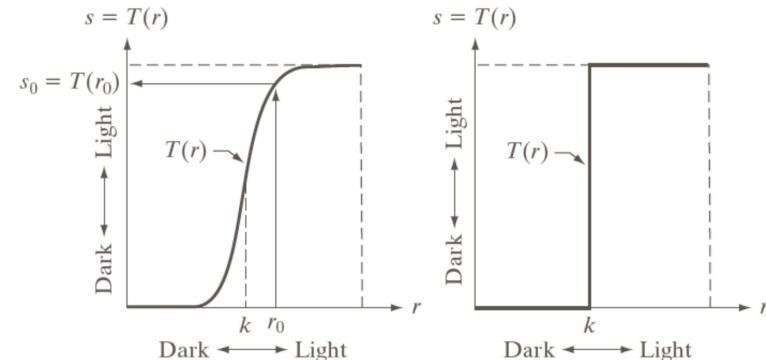
$(N_1 \times N_2 \times \dots)$: pixel density=**resolution**

c : RGB channels for any **color**

**What kind of operations can we do on
digital images?**

Intensity transformations

Transform each pixel value pointwise.



LearnOpenCV.com

Pixel level operations

- 0D linear function (pixel) $ax + b$

90	3	241	65	213
242	143	122	71	8
203	103	45	167	240
29	239	2	93	218
243	242	183	54	123

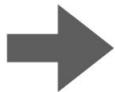
- 1D linear function (color)

$$\begin{bmatrix} 0.19 & 0.77 & 0.39 \\ 0.17 & 0.69 & 0.35 \\ 0.13 & 0.53 & 0.27 \end{bmatrix} \begin{bmatrix} x_R \\ x_G \\ x_B \end{bmatrix}$$

- Any function (transfer function)

Input	Output
0	0
...	...
100	50
...	...
255	127

Brightness adjustment = addition / subtraction



Pixel arithmetic

243	222	171
203	145	137
209	132	193

123	102	51
83	25	17
89	12	73

$$p \rightarrow p - 120$$

Input pixel value

Output pixel value

```
I = imread('image.png')  
I = I - 120
```



Caution: numerical overflow

$$100 + 200 = ?$$

Overflow: $300 - 256 = \mathbf{44}$

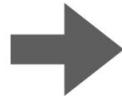


Uint8 image arithmetic

- convert to “Float”
- do arithmetic
- clip to [0,255]
- convert it back to “uint8”

```
np.clip((im_shop.astype(float)+60),0,255).astype(np.uint8)
```

Contrast adjustment = multiplication/division



We want bright to be brighter, dark to be darker.

If we only multiplied, this wouldn't work.

121	111	85
101	72	68
104	66	96

181	166	127
151	108	102
156	99	144

 p  $p * 1.5$

Input pixel value

Output pixel value

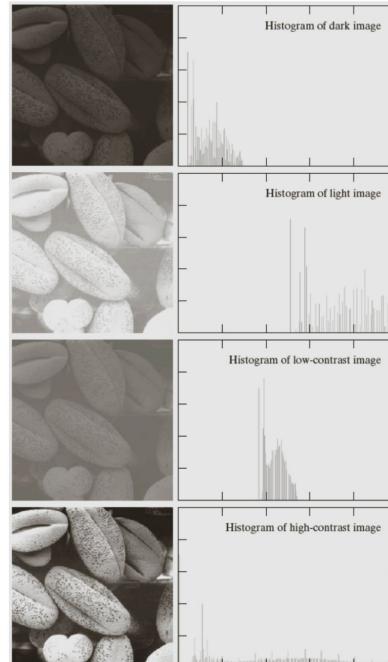


Bright gets brighter, dark **also** gets brighter.

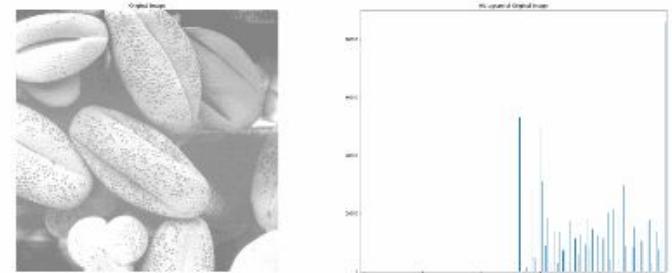
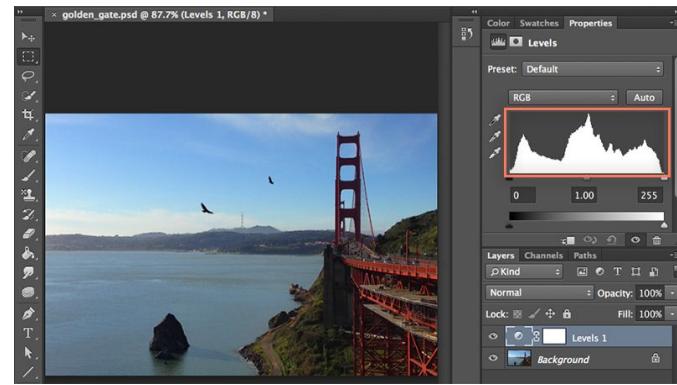


Think of “brightness” vs “darkness” as a distribution of pixel values = a histogram.

- $h(r_i) = n_i$
 - Histogram: number of times intensity level r_i appears in the image
- $p(r_i) = n_i/NM$
 - normalized histogram
 - also a probability of occurrence

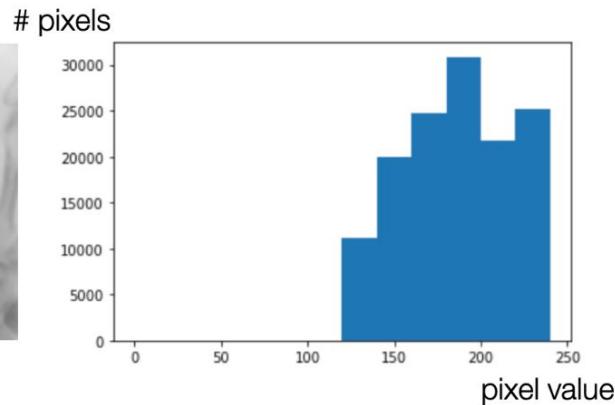


imgflip.com



What are the operations to improve contrast?

I : (min, max) = 120, 230

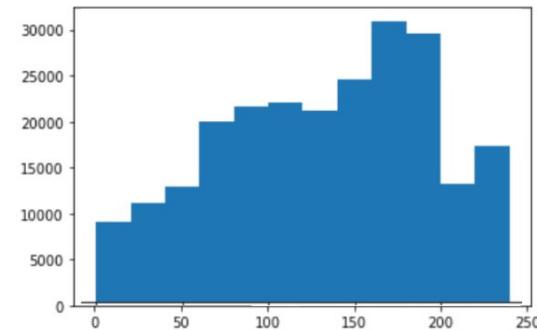


Pixel value histogram

?



I : (min, max) = 0, 255

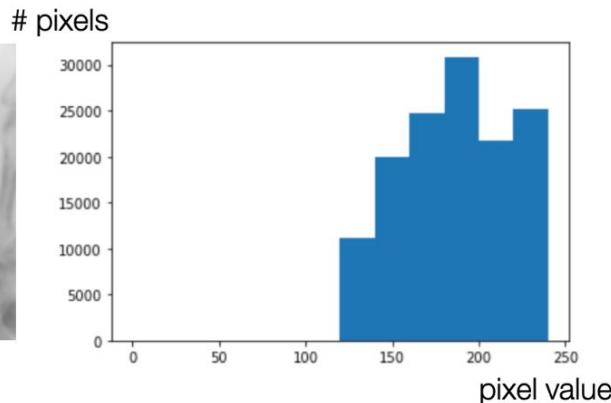


(min, max)

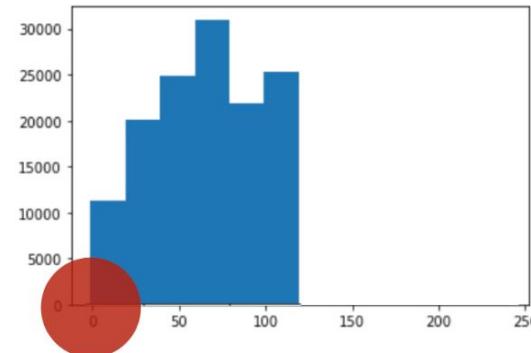
I : (min, max) = 120, 230

$I = I - I.\text{min}()$

0, 110

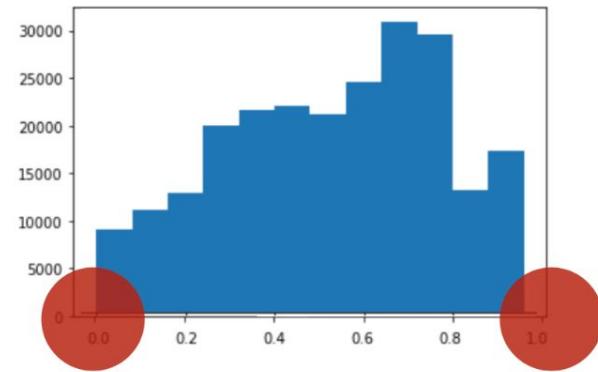
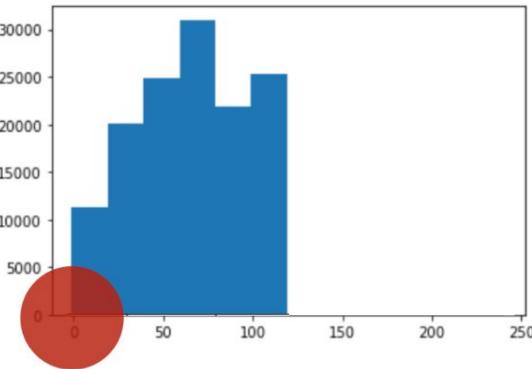


Pixel value histogram



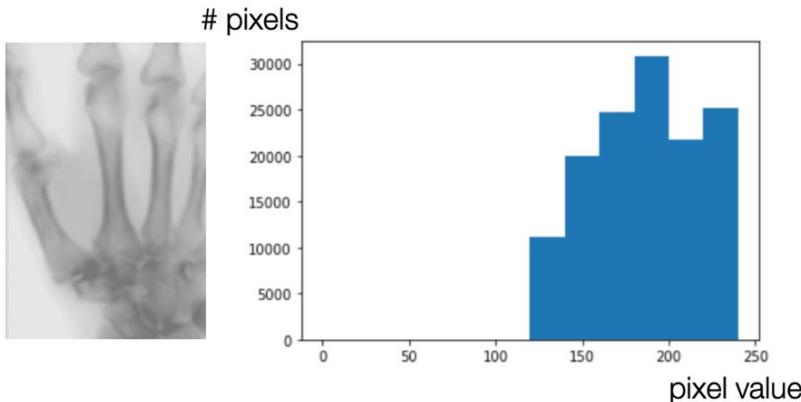
Step 2: scale to right endpoint (min, max)

$I = I - I.\text{min}()$ 0, 110
 $I = I / (I.\text{max}() - I.\text{min}())$ 0, 1



(min, max)

I : (min, max) = 120, 230



Pixel value histogram

$I = I - 120$

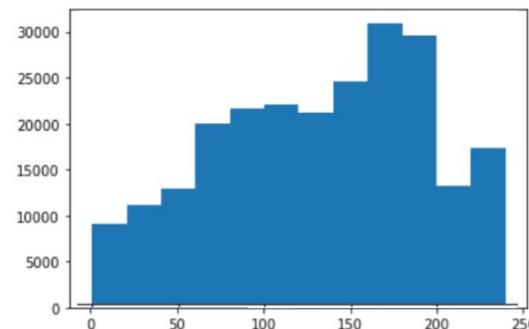
0, 110

$I = I / (I.\max() - I.\min())$

0, 1

$I = I * 255$

0, 255



A combination of multiplicative and additive operations needed.

Intensity vs. color transformations

Intensity operations are arithmetic - pixelwise.

Color transformations require matrix operations since operating on a vector.

1 pixel



Intensity
(scalar: arithmetic)



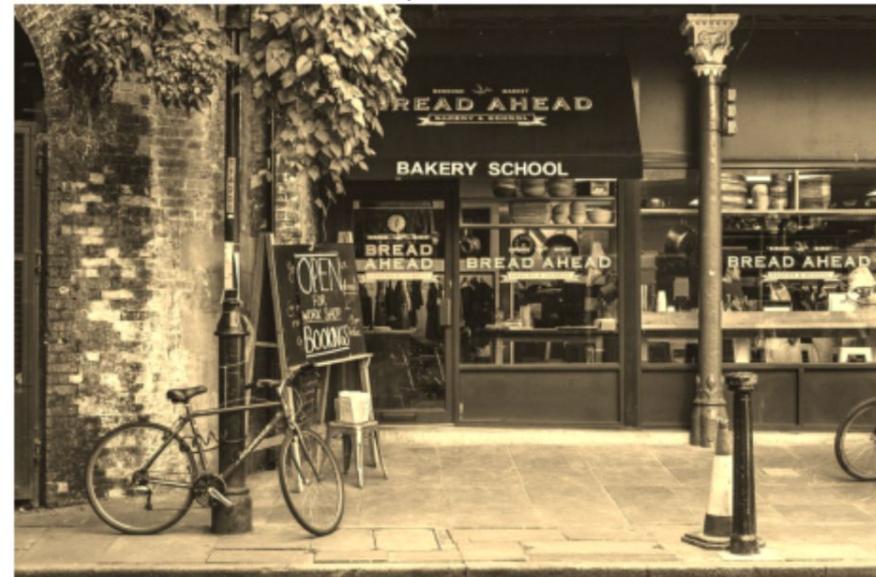
RGB
(vector: linear algebra)

Color filter

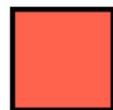
input



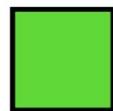
sepia effect



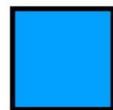
Matrix operation on each pixel color vector



$$= 0.19 R + 0.77 G + 0.39 B$$



$$= 0.17 R + 0.69 G + 0.35 B$$



$$= 0.13 R + 0.53 G + 0.27 B$$

RGB_new

(Image range: 0-1)

$$= \begin{bmatrix} 0.19 & 0.77 & 0.39 \\ 0.17 & 0.69 & 0.35 \\ 0.13 & 0.53 & 0.27 \end{bmatrix} \begin{bmatrix} R \\ G \\ B \end{bmatrix}$$

Matrix operation on a tensor

Color images are $N \times M \times C$ tensors.

“Vectorize” each image into a $NM \times C$ matrix.

Each entry is a pixel.

Transform each pixel and then reshape it back to the $N \times M \times C$ form.

```
color_trans = np.array([[0.189, 0.168, 0.131],  
                      [0.769, 0.686, 0.534],  
                      [0.393, 0.349, 0.272]])
```

```
# reshape image into Nx3  
im_shop_reshape = im_shop.reshape(-1,3)
```

```
# convert image range into 0-1  
im_shop_reshape = im_shop_reshape/ 255.0
```

```
# Sepia color transform  
im_shop_sepia = np.matmul(im_shop_reshape, color_trans)
```

```
# reshape it back  
im_shop_sepia = im_shop_sepia.reshape(im_shop.shape)
```



Non-linear color transformations

input



summer effect



Lookup tables

Transfer function

$$f: p \rightarrow p'$$

Input pixel Output pixel

0	0
...	...
100	50
...	...
255	127

Example:
 $f_1(p) = p/2$

Key observation:

Array index -> value

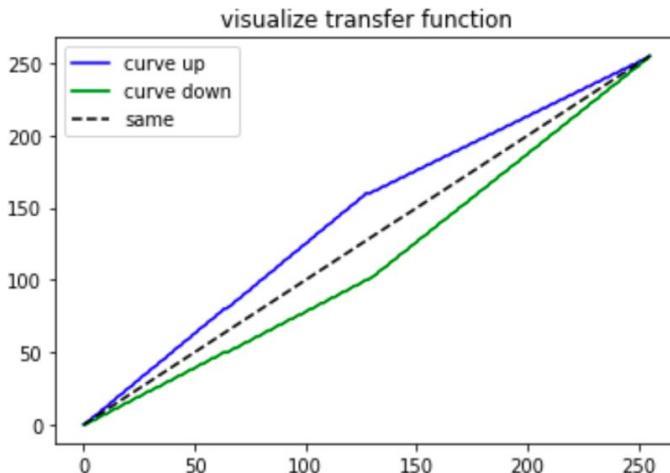
```
# load in a matrix (uint8: 0-255)
I = imread('image.png')
```

```
# transfer function
f1 = np.arange(255) // 2
```

```
# apply point-level operation
# (relabel input pixel)
I_out = f1[I]
```



Example of non-linear transfer function



Goal: more red and less blue

1. Create the transfer function

```
transfer_up = np.zeros(256, np.uint8)
transfer_up[:64] = np.linspace(0, 80, 64)
transfer_up[64:128] = np.linspace(80, 160, 64)
transfer_up[128:256] = np.linspace(160, 255, 128)
```

2. Apply the transfer function

```
im_shop_summer = im_shop.copy()
im_shop_summer[:, :, 0] = transfer_up[im_shop_summer[:, :, 0]]
im_shop_summer[:, :, 2] = transfer_down[im_shop_summer[:, :, 2]]
```

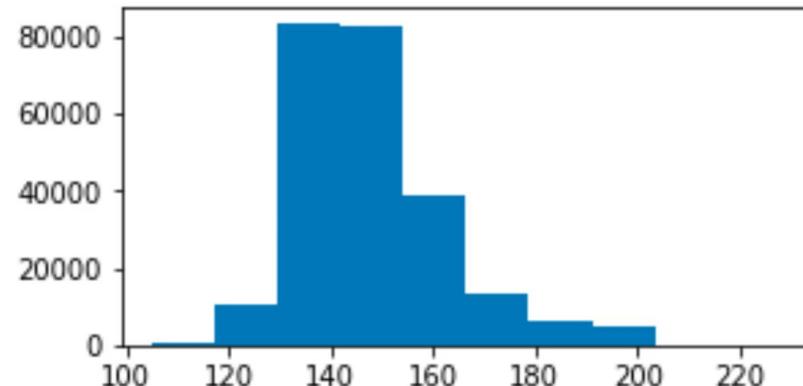


Even a more non-linear transformation needed

- Too many gray pixels → not much contrast



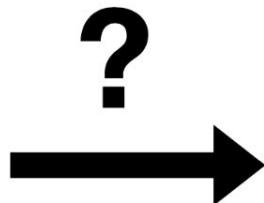
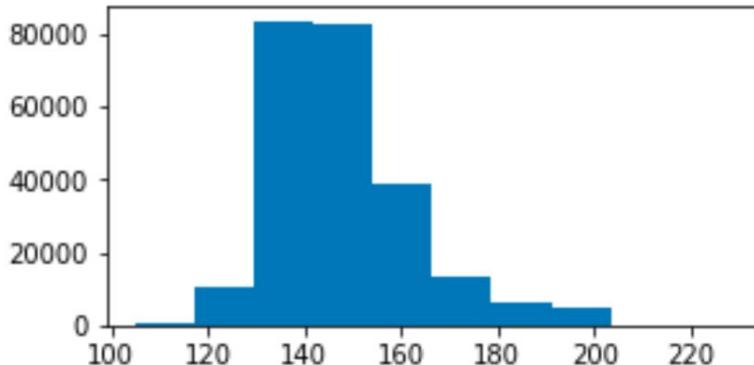
Histogram of pixel values



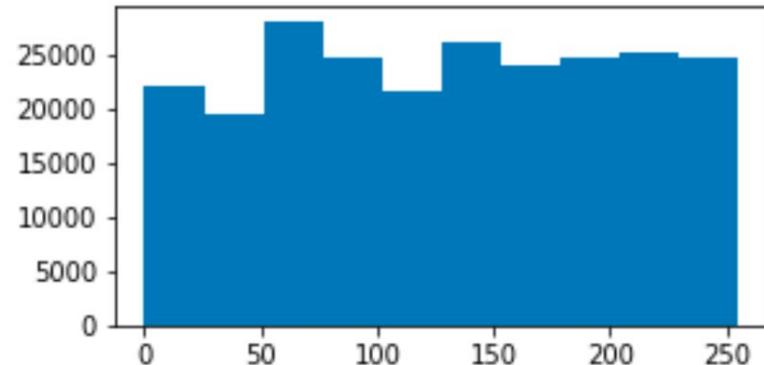
Would additive or subtractive functions work?

We want to cover all the brightness levels

Input pixel histogram



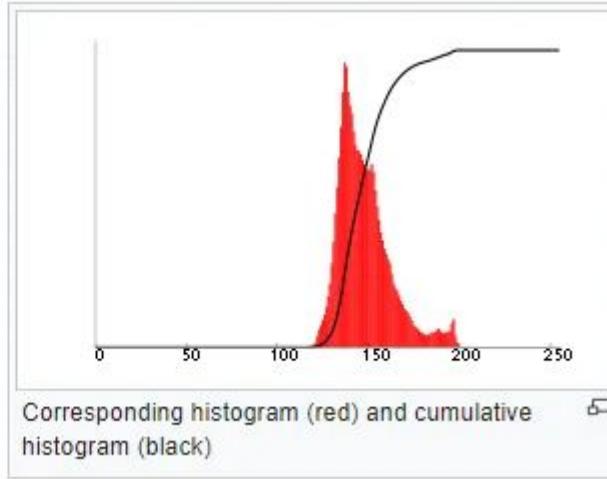
Output pixel histogram



We can design a lookup table to achieve this.



Cumulative distribution function (CDF)



For each value in a set x , CDF computes the probability that x is greater than all other values.

We can use this as a lookup table (with some transformations) to uniformize the histogram.

Image



pixel value count

pixel value	count
50	10
100	20
255	5



(N pixels in total)

?? pixels with value $\leq K$

pixel value cumulative

50	10
100	30
255	35

`im_desert_cdf = np.cumsum(unique_count)`

`unique_id, unique_count = np.unique(im_desert, return_counts=True)`



pixel value New value

50	0
100	??
255	255

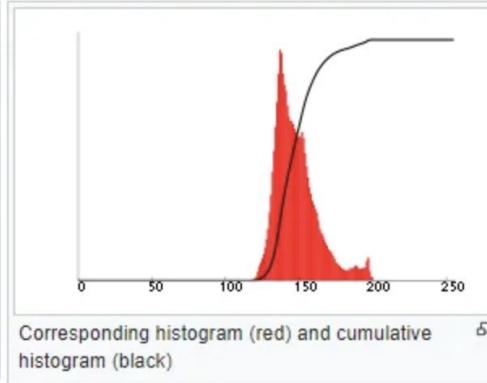
$$(30-10)/(35-10)*255=219$$

$$(x - \text{cdf}_{min}) * \frac{255}{N - \text{cdf}_{min}}$$





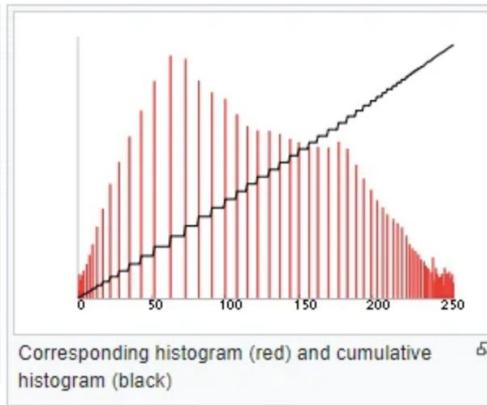
Before Histogram Equalization



Corresponding histogram (red) and cumulative histogram (black)



After Histogram Equalization



Corresponding histogram (red) and cumulative histogram (black)

We covered pixel level operations - any questions?

- 0D linear function (pixel) $ax + b$

90	3	241	65	213
242	143	122	71	8
203	103	45	167	240
29	239	2	93	218
243	242	183	54	123

- 1D linear function (color)

$$\begin{bmatrix} 0.19 & 0.77 & 0.39 \\ 0.17 & 0.69 & 0.35 \\ 0.13 & 0.53 & 0.27 \end{bmatrix} \begin{bmatrix} x_R \\ x_G \\ x_B \end{bmatrix}$$

- Any function (transfer function)

Input	Output
0	0
...	...
100	50
...	...
255	127

When pixel operations are not enough?



Clean image



Real-world image

Is there any pixel operation that can “fix” this image?

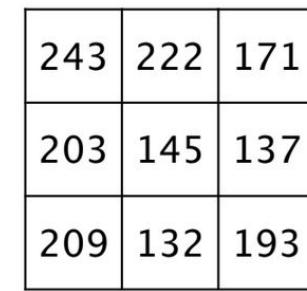
Need to widen the input of our transformation functions



$$f: p \rightarrow p'$$

Input pixel

Output pixel



$$f: pa \rightarrow p'$$

Input **patch**

Output pixel

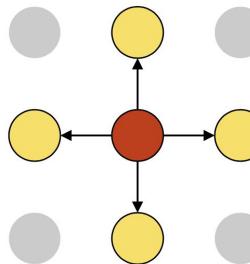
Brief detour on neighborhoods and patches

To denote patches - we need to denote “neighborhoods”:

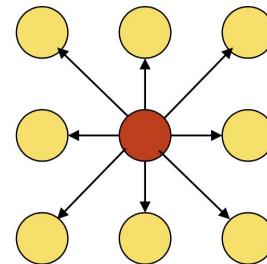
1. Impose topological structure e.g.. “Pixels lie on a grid graph”
 - a. 4 or 8 cell topologies.
2. Define adjacency
 - a. Adjacency is the set of nodes that are connected.
3. Define a path
 - a. Path is a series of pixels that are sequentially adjacent.
4. Define distance
 - a. The distance between two pixels is the length of the shortest path.
 - b. Ex. D4 or D8 distances.

A patch of radius r around a pixel is the set of all of pixels that are less than a r -distance away.

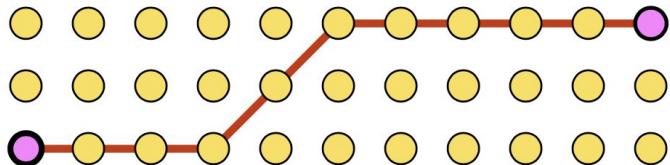
4 connected



8 connected

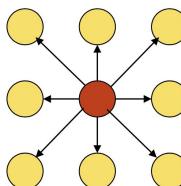


D8 distance



D8 patch of radius 2

8 connected



2	2	2	2	2
2	1	1	1	2
2	1	0	1	2
2	1	1	1	2
2	2	2	2	2

Example: median filter on patch size 1 (in d8)

243	222	171
203	145	137
209	132	193



121	111	85
101	183	68
104	66	96

$$f: pa \rightarrow p'$$

Input **patch**

Output pixel

What's the output?

30	31	32	3	4
0	6	99	30	90
99	35	33	32	98
0	90	90	36	31
32	31	0	90	90

Input

				?

Output

$$p' = \text{median}(p)$$



30	31	32	3	4
0	6	99	30	90
99	35	33	32	98
0	90	90	36	31
32	31	0	90	90

Input

Output



30	31	32	3	4
0	6	99	30	90
99	35	33	32	98
0	90	90	36	31
32	31	0	90	90

Input

			32	32

Output

30	31	32	3	4
0	6	99	30	90
99	35	33	32	98
0	90	90	36	31
32	31	0	90	90

Input

		32	32	32
		35	35	36
		33	35	36

Output

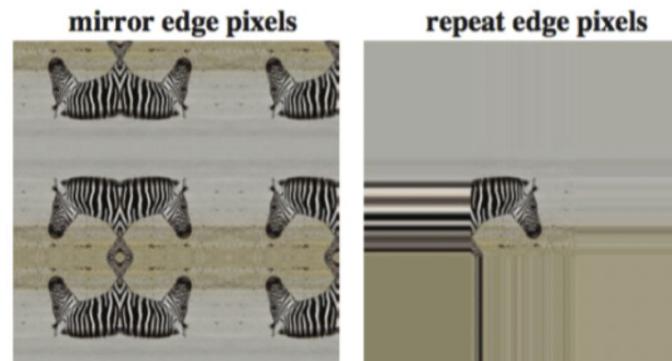
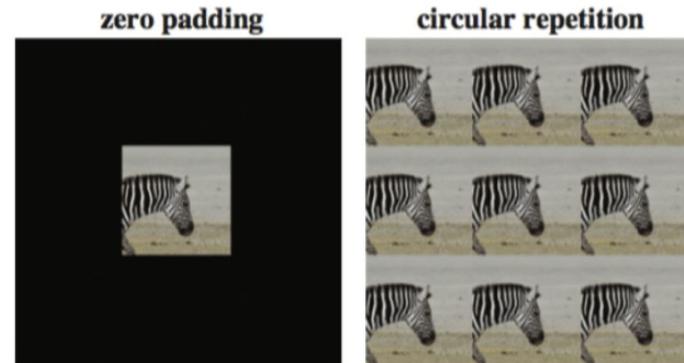
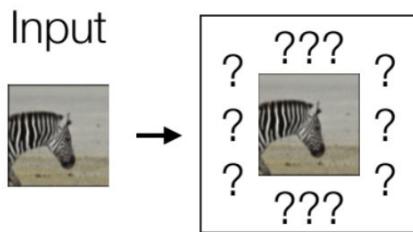
Brief issue detour: patch overflows out of image

0	0	0	0	0	0	0
0	90	90	90	90	0	0
0	90	90	90	90	0	0
0	90	90	90	90	0	0
0	90	0	90	90	0	0
0	90	90	90	90	0	0
0	0	0	0	0	0	0

	23	54	54	23	20	
	60	90	54	23	20	
	50	80	80	54	30	
	?	50	80	80	54	30
	30	50	50	23	20	



Solution: padding



Median filter is a non-parametric filter but it's non-linear



Ideal image



Real-world image

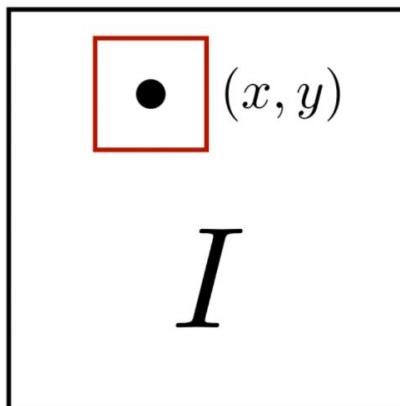


Median filter result

What are some other examples of filters?



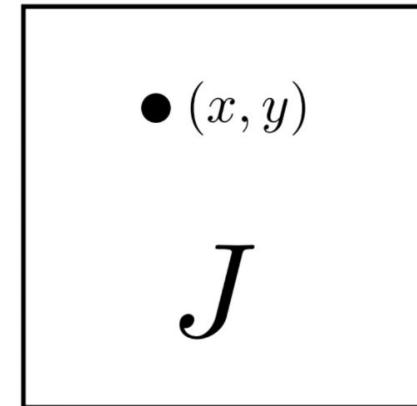
Patch centered at (x, y)



Input image

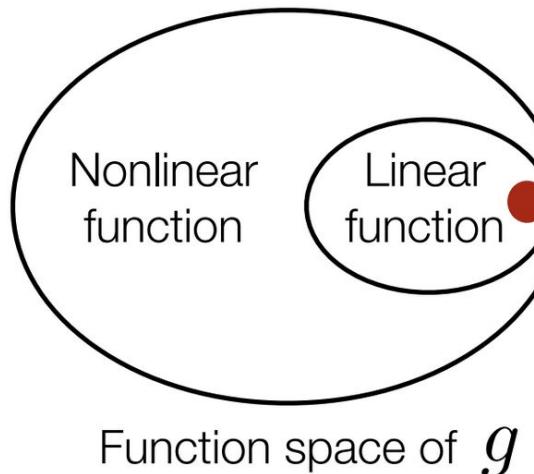
Filter
 g

Image intensity at (x, y)



Output filtered image

Linear vs. non-linear filters



Linear +
Shift-invariant

- Linear Function

$$g(\alpha I_1 + \beta I_2) = \alpha g(I_1) + \beta g(I_2)$$

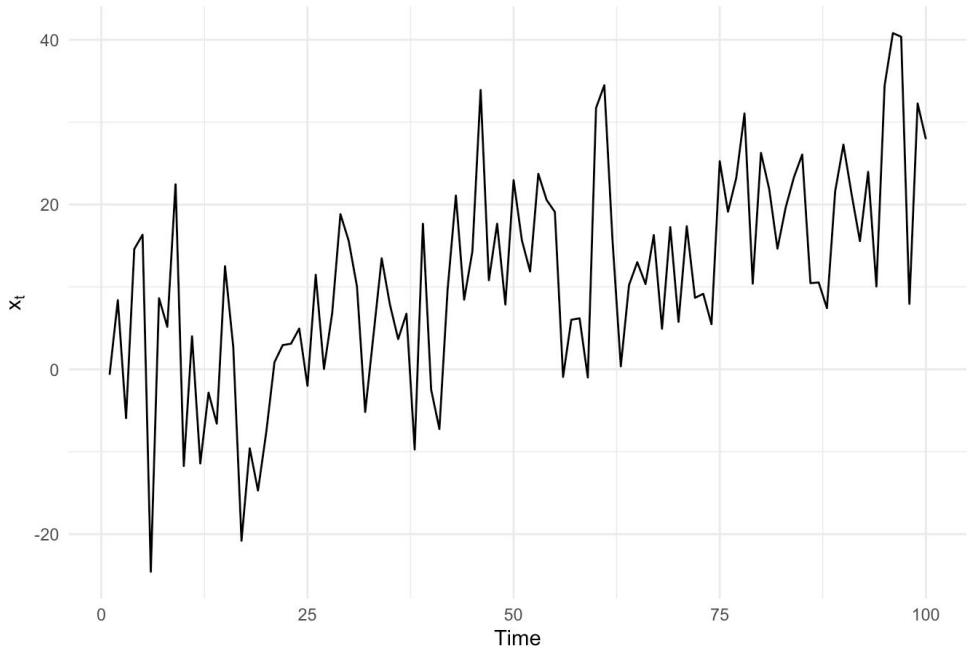
- Shift invariance

Same function g for all image patches



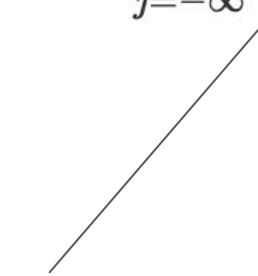
1D Linear filtering example

Input x



Filtered 1D data

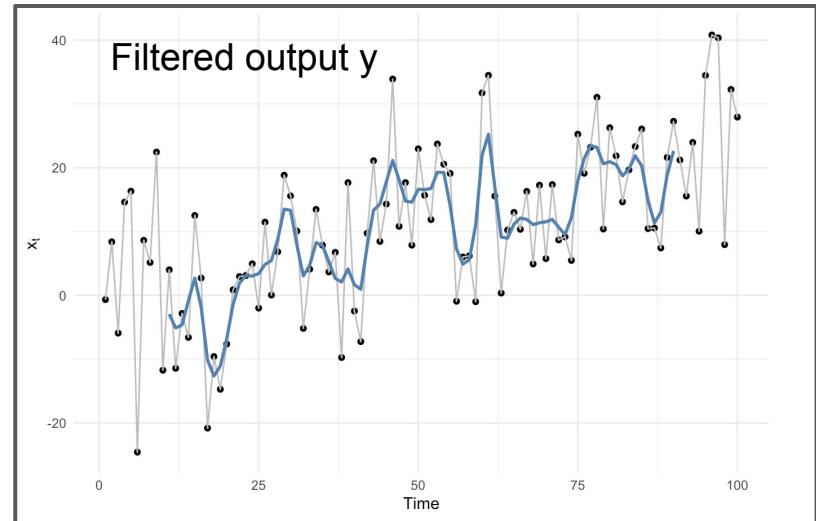
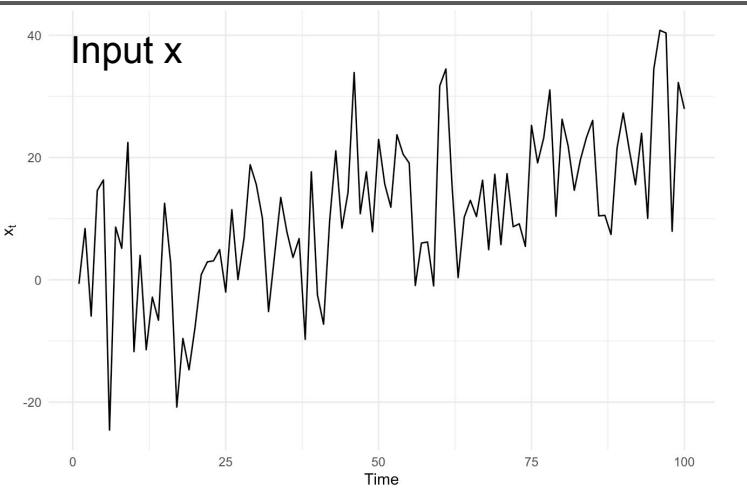
$$y_t = \sum_{j=-\infty}^{\infty} \beta_j x_{t-j}$$



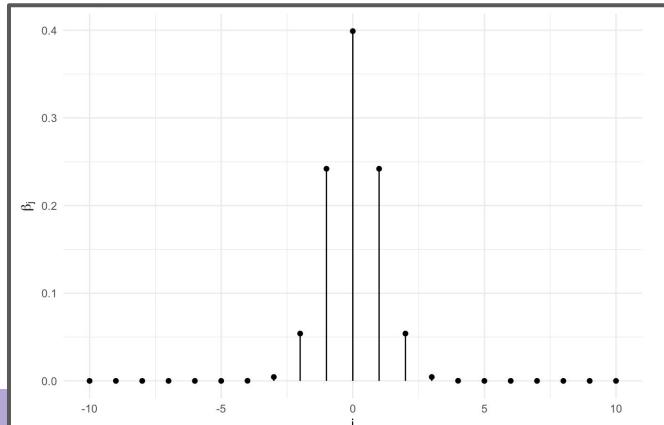
β_j : filter parameters



Example: Low pass filter

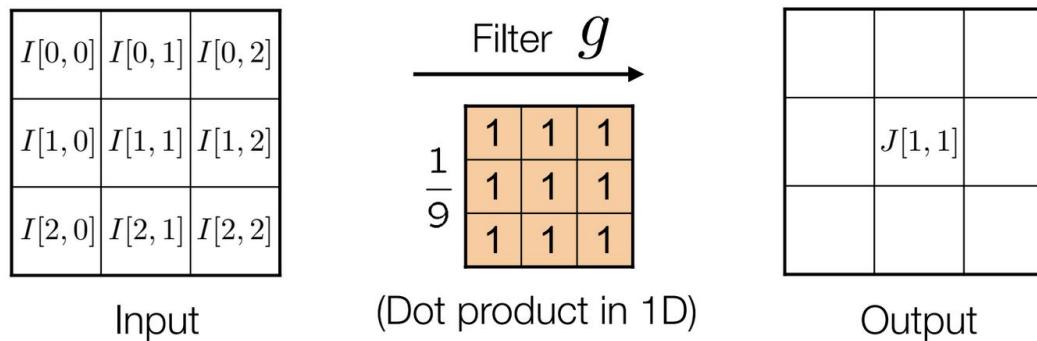


B_j : filter parameters



Simple linear box filter in 2D (moving average)

$I[x, y]$: Image intensity at pixel (x,y)



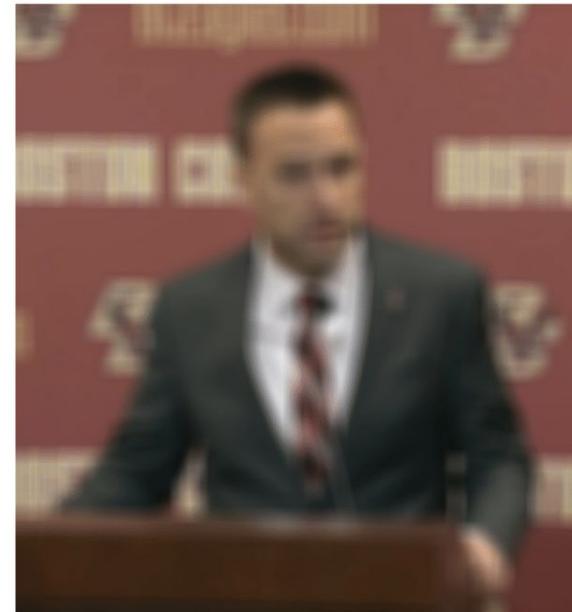
$$J[1, 1] = \frac{1}{N} \sum_{x,y} I[x, y]$$



Box filter: blur effect



Input Image



Output: Filtered image
(mimic focus blur)

Basic filters can help computers “perceive” edges, corners and denoise errors

$$\frac{1}{K^2} \begin{bmatrix} 1 & 1 & \cdots & 1 \\ 1 & 1 & \cdots & 1 \\ \vdots & \vdots & 1 & \vdots \\ 1 & 1 & \cdots & 1 \end{bmatrix}$$

$$\frac{1}{16} \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix}$$

$$\frac{1}{256} \begin{bmatrix} 1 & 4 & 6 & 4 & 1 \\ 4 & 16 & 24 & 16 & 4 \\ 6 & 24 & 36 & 24 & 6 \\ 4 & 16 & 24 & 16 & 4 \\ 1 & 4 & 6 & 4 & 1 \end{bmatrix}$$

$$\frac{1}{8} \begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}$$

$$\frac{1}{4} \begin{bmatrix} 1 & -2 & 1 \\ -2 & 4 & -2 \\ 1 & -2 & 1 \end{bmatrix}$$

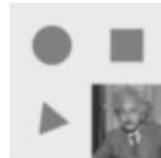
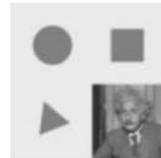
$$\frac{1}{K} \begin{bmatrix} 1 & 1 & \cdots & 1 \end{bmatrix}$$

$$\frac{1}{4} \begin{bmatrix} 1 & 2 & 1 \end{bmatrix}$$

$$\frac{1}{16} \begin{bmatrix} 1 & 4 & 6 & 4 & 1 \end{bmatrix}$$

$$\frac{1}{2} \begin{bmatrix} -1 & 0 & 1 \end{bmatrix}$$

$$\frac{1}{2} \begin{bmatrix} 1 & -2 & 1 \end{bmatrix}$$



(a) box, $K = 5$

(b) bilinear

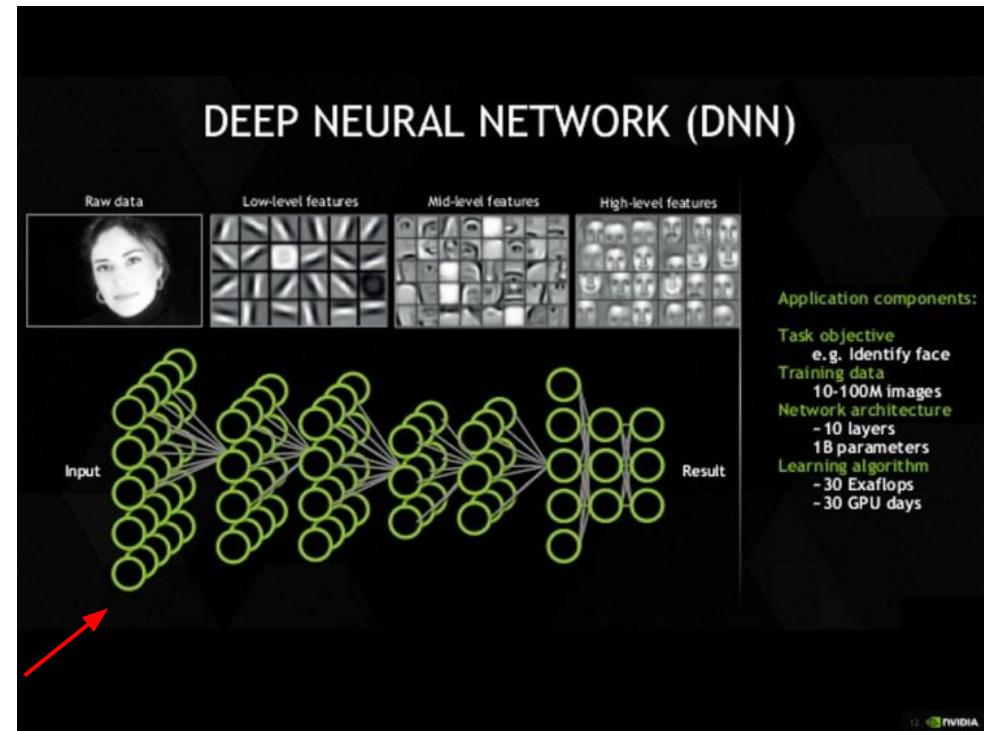
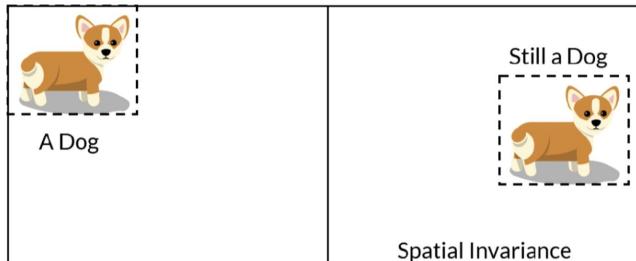
(c) “Gaussian”

(d) Sobel

(e) corner

More non-linear filters are required for higher order recognition

Deep learning models (built upon filters) need to be shift-invariant.



We are here



Next: code lab

Next lecture - more filtering + convolutions + geometric transformations

